

Hadoop Online Quiz

Following quiz provides Multiple Choice Questions (MCQs) related to **Hadoop Framework**. You will have to read all the given answers and click over the correct answer. If you are not sure about the answer then you can check the answer using **Show Answer** button. You can use **Next Quiz** button to check new set of questions in the quiz.



Q 1 - HDFS block size is larger as compared to the size of the disk blocks so that

- A - Only HDFS files can be stored in the disk used.
- B - The seek time is maximum
- C - Transfer of a large files made of multiple disk blocks is not possible.
- D - A single file larger than the disk size can be stored across many disks in the cluster.

Answer : D

Hide Answer

Q 2 - Zookeeper ensures that

- A - All the namenodes are actively serving the client requests
- B - Only one namenode is actively serving the client requests

C - A failover is triggered when any of the datanode fails.

D - A failover can not be started by hadoop administrator.

Answer : B

Hide Answer

Q 3 - For a HDFS directory the replication factor(RF) is

A - same as the RF of the files in that directory

B - Zero

C - 3

D - Does not apply.

Answer : D

Show Answer

Q 4 - The hadfs command put is used to

A - Copy files from local file system to HDFS.

B - Copy files or directories from local file system to HDFS.

C - Copy files from from HDFS to local filesystem.

D - Copy files or directories from HDFS to local filesystem.

Answer : B

Hide Answer

Q 5 - Which of the below apache system deals with ingesting streaming data to

hadoop

A - Ozie

B - Kafka

C - Flume

D - Hive

Answer : C

Hide Answer

Q 6 - YARN stands for

A - Yahoo's another resource name

B - Yet another resource negotiator

C - Yahoo's archived Resource names

D - Yet another resource need.

Answer : B

Hide Answer

Q 7 - When you increase the number of files stored in HDFS, The memory required by namenode

A - Increases

B - Decreases

C - Remains unchanged

D - May increase or decrease

Answer : A

Hide Answer

Q 8 - In a HDFS system with block size 64MB we store a file which is less than 64MB. Which of the following is true?

- A - The file will consume 64MB
- B - The file will consume more than 64MB
- C - The file will consume less than 64MB.**
- D - Can not be predicted.

Answer : C

Hide Answer

Q 9 - What is distributed cache?

- A - The distributed cache is special component on name node that will cache frequently used data for faster client response. It is used during reduce step.
- B - The distributed cache is special component on data node that will cache frequently used data for faster client response. It is used during map step.**
- C - The distributed cache is a component that caches java objects.
- D - The distributed cache is a component that allows developers to deploy jars for Map-Reduce processing.

Answer : B

Hide Answer

Q 10 - In order to apply a combiner, what is one property that has to be satisfied by the values emitted from the mapper?

A - Combiner can be applied always to any data

B - Output of the mapper and output of the combiner has to be same key value pair and they can be heterogeneous

C - Output of the mapper and output of the combiner has to be same key value pair. Only if the values satisfy associative and commutative property it can be done.

Answer : C

Hide Answer

New Quiz